

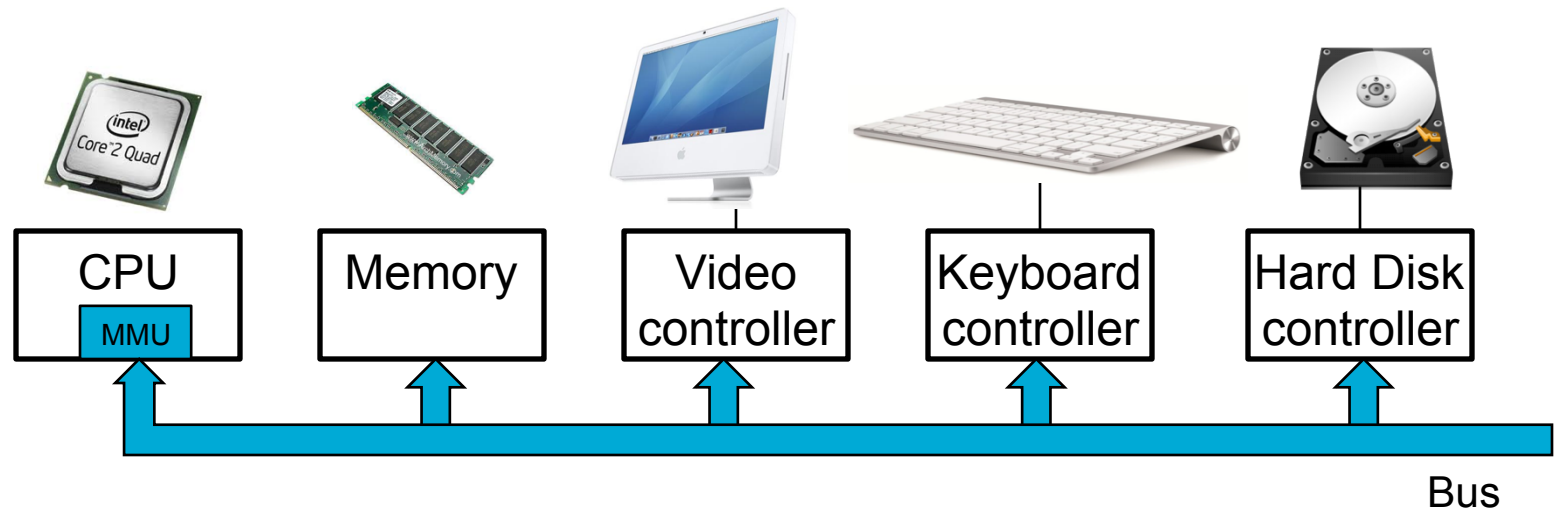
Architecture and OS

To do ...

- ❑ Architecture impact on OS
- ❑ OS impact on architecture
- ❑ Next time: OS components and structure


Computer architecture and OS

- OS is intimately tied to the hardware it runs on
 - OS design is impacted by it
 - OS needs result on new architectural features
- Abstract model of a simple computer



Processor



- The brain with a basic operation cycle
 - Fetch next instruction
 - Decode it to determine type & operands
 - Execute it
- ... and a specific set of instructions
 - Architecture specific – x86 != ARM
- Plus registers, since memory is slow
 - General registers and special ones (PC, SP)
- Clearly simplistic model
 - Pipelining, superscalar, multicore



Memory


- Ideal – fast, large, cheap, persistent
- Real – storage hierarchy
 - Registers
 - Internal to the CPU & just as fast
 - Cache
 - If word needed is in cache, get it in ~ 2 cycles
 - Main memory $\sim 100\times$ slower
 - Hard disk – from nsecs to msec (moving the arm 5-10msec)
 - Magnetic tape
 - Cloud storage

First core-based memory: IBM 405 Alphabetical Accounting Machine



Architectural trends impact OS design ...

- Processing power
 - Doubling every 18 months (100x per decade)
 - but power is a serious issue
- Disk capacity
 - Double every 12 months (1000x per decade)
 - 1961 ~ \$4,440/MB
 - 2014 ~ \$0.00004/MB



1961 IBM 1301
~26MB ~\$115,500
(~890k in 2013)

Architectural trends impact OS design

- Memory

- Same and for the same reason
 - 1980 64KB \$405.00 (\$6,480/MB) ...
 - 2014 4GB \$29.99 (\$0.007/MB)*

- Optical bandwidth today

- 10x as fast as disk capacity
- 100x as fast as processor performance
- Doubling every 9 months (Butter's law)
- Latency is the limiting factor, but there's room
 - Time to get the HTML index page from popular sites is, in the median, $34 * c\text{-latency}$ (light's shortest path rtt)

... and OS needs shape the architecture

- Arch support can simplify/complicate OS tasks
 - E.g., early PC OS (DOS) lacked support for virtual memory, partly because HW lacked key features
- Features built primarily to support OS's
 - Protected modes of execution (kernel vs. user)
 - System calls (and software interrupts)
 - Memory protection
 - I/O control operations
 - Timer (clock) operation
 - Interrupts and exceptions
 - Synchronization instructions

Consider timesharing

- Multiprogramming & timesharing are useful
 - Multiprogramming – “using different parts of the hardware at the same time for different tasks”
 - Timesharing – “several persons making use of the computer at the same time”
- but
 - How to protect programs from each other and the kernel from all?
 - How to handle relocation? OS may need to run a given program at != times from != locations

For protection

- Restrict some instructions to the OS
 - e.g. Directly access I/O devices
- How does the CPU know if a protected instructions should be executed?
 - Architecture must support 2+ modes of operation
 - Mode is set by status bit in a protected register
 - User programs execute in user mode, OS in kernel mode
- Protected instructions can only execute in kernel mode

Crossing protection boundaries

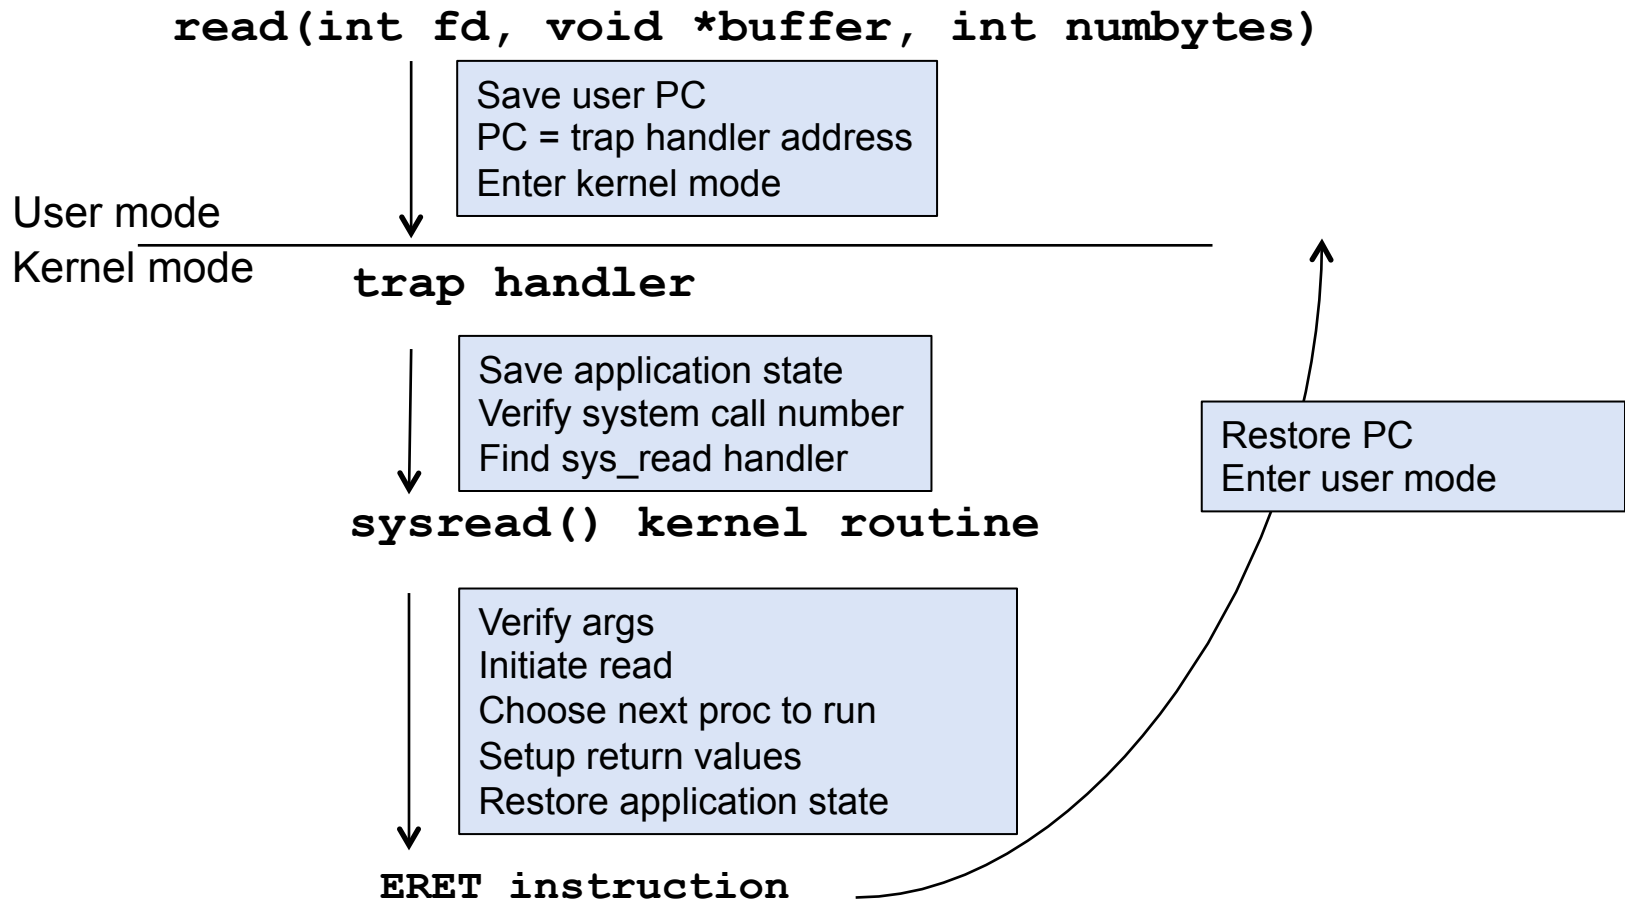
- How can apps do something privileged?
 - e.g. How do you save a file if you can't do I/O?
- User programs must call an OS procedure
 - Ask the OS to do it
 - OS defines a set of **system calls**
 - User-mode program makes a system call
 - How does the user to kernel-mode transition happen?

Crossing protection boundaries

- The system call ...
 - Causes an exception which vector to a kernel handler
 - Passes a parameter indicating which syscall is
 - Saves caller's state so it can be restored
 - *What would happen if the kernel didn't save state?*
 - OS must verify caller's parameters
 - *Why should it do that?*
 - Must have a way to go back to user once done
 - A special instruction sets PC to the return address and the execution mode to user

Crossing protection boundaries

- A system call



- *A bit like a regular subroutine call, right?*

Exception handling and protection


- All entries to the OS use the same mechanism
 - Acquiring privileged mode and branching to trap handler are inseparable
- Interrupts, exceptions and traps
 - Interrupt: asynchronous, caused by an external HW event
 - Exception: synchronous; unexpected, automatically generated (coerced, no requested); divide by zero
 - Trap: synchronous; programmer initiated, expected transfer of control to a special handler

Next issue – Memory relocation

- Simplest model – base + limit
 - Base (start) of program + limit registers
 - Used by CDC 6600 (supercomputer) and Intel 8088
 - Changing program means changing base+limit
 - Solves relocation and protection (on CDC 6600)
 - Cost 2 registers + cycle time incr

- More sophisticated alternatives
 - 2 base and 2 limit registers for text & data; allow sharing program text
 - Paging, segmentation, virtual memory

If \geq base and $<$ base + limit, OK
else trap to OS with error



OS needs shape the architecture – I/O

- I/O Device

- Device + Controller (simpler I/F to OS; think SCSI)
 - Read sector x from disk y → (disk, cylinder, sector, head), ...

- How does the kernel start an I/O?

- Special I/O instructions
- Memory-mapped I/O

- How does it notice when the I/O is done?

- Polling – *are we done yet?*
- Interrupts – *let me know when you are done?*

- How does it exchange data with I/O device?

- Programmed I/O
- Direct Memory Access (DMA)

OS control flow

- OSs are event driven
 - Once booted, all entry to kernel happens as result of an event (e.g. signal by an interrupt), which
 - Immediately stops current execution
 - Changes to kernel mode, event handler is called
- Kernel defines handlers per event type
 - Specific types are defined by the architecture
 - e.g. timer event, I/O interrupt, system call trap

Timers

- How can the OS retain control when a program gets stuck in an infinite loop?
 - HW timer that generates a periodic interrupt
 - Before it transfers to a user program, OS loads timer with a time to interrupt
 - When time's up, interrupt transfers control back to OS
 - OS pick a program to schedule next (which one?)
- Should the timer be privileged?
 - For reading or for writing?

Synchronization

- Issues with interrupts
 - Executing code can interfere with interrupted code
 - OS must be able to synchronize concurrent processes
- Synchronization
 - Guarantee that short instruction sequences (e.g. read-modify-write) execute atomically
 - Two methods
 - Turn off interrupts, execute sequence, re-enable interrupts
 - Have special, complex atomic instructions – test-and-set

Management of concurrency & asynchronous events is the biggest difference between systems-level & traditional app programming

Summary & preview

- This is far from over – new architectural features are still being introduced
 - Support for virtual machine monitors
 - Hardware transaction support
 - Support for security
 - Low latency persistent memory
 - ...
- Transistors are free so Intel/AMD/... need to find applications that require new hardware that you would want to buy ...