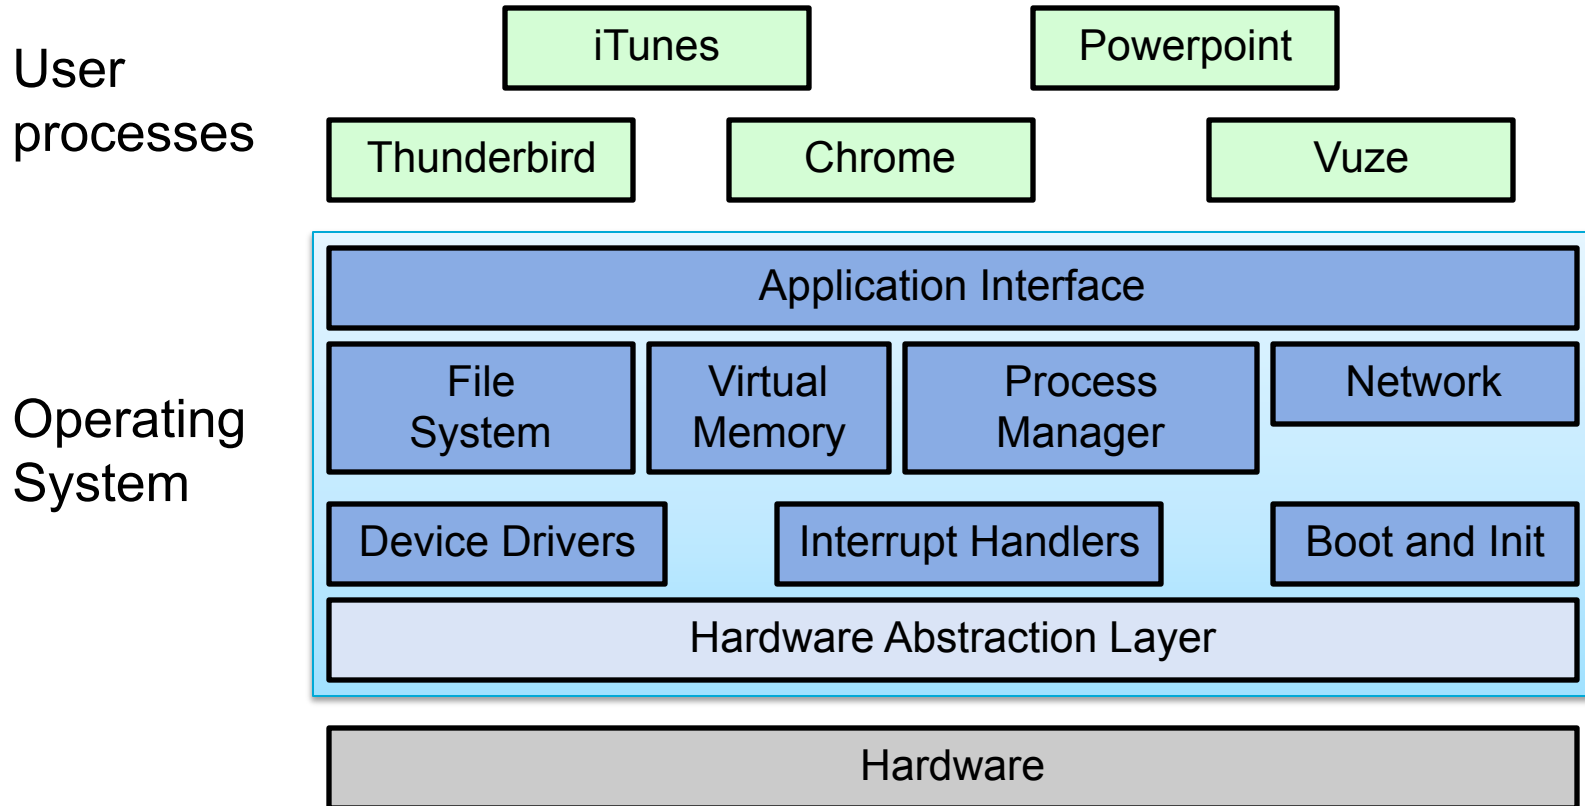


OS concepts and structure

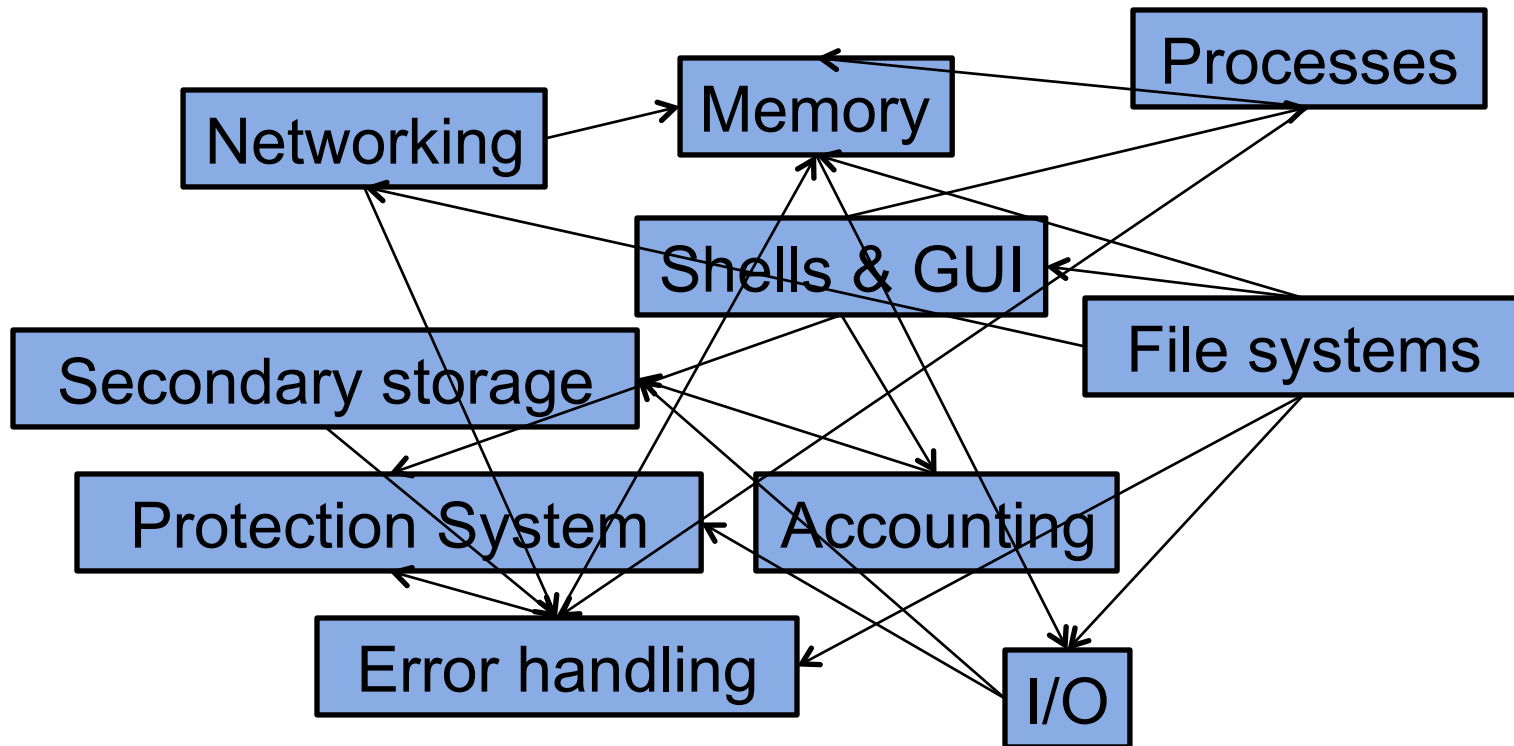
To do ...

- ❑ OS components & interconnects
- ❑ Structuring OSs
- ❑ Next time: Processes

Between hardware and your apps



Major OS components & abstractions



OS Views

- Perspectives, “OS as ...
 - ... the services it provides
 - ... its components and interactions
- Services to ...
 - Users via a GUI or a command interpreter/shell
 - Programmers via system calls
 - Some services are for convenience: FS management
 - Some to ensure efficient operation: Resource allocation

GUI or command interpreter (shell)

- GUI

- Friendlier (desktop), if sometimes limiting
- Xerox PARK Alto >> Apple >> Windows >> Linux

- Command interpreter

- Handle (interpret and execute) user commands
- Could be part of the OS: MS DOS, Apple II
- Or just a special program: UNIX, Win XP
- The command interpreter could
 - Implement all commands
 - Simply understand what program to invoke and how (UNIX)

System calls

- Low-level interface to services for applications
- Higher-level requests get translated into sequence of system calls
- Writing `cp` – copy source to destination
 - Get file names
 - Open source
 - Create destination
 - Loop
 - Read from source
 - Copy to destination
 - Close destination
 - Report completion
 - Terminate

Major OS components & abstractions

- Processes
- Memory
- I/O
- Secondary storage
- File systems
- Protection
- Accounting
- Shells & GUI
- Networking

Processes

- An OS executes many kind of activities
 - Each encapsulated in a process
- A program in execution
 - Address space, set of registers, OS resources
 - Threads and processes – for now consider each process to have a single thread (we'll revisit this later)
- To get a better sense of it
 - What data do you need to re-start a suspended process?
 - Where do you keep this data?
 - What is the process abstraction interface offered by the OS?

Memory management

- Main memory – directly accessed for CPU
 - Programs must be in memory to execute
 - Memory access is fast (e.g., 60 ns to load/store), but not persistent (won't survive power failures)
- OS must
 - Allocate memory space to processes
 - Decide how to allocate it to each process
 - Deallocate it when needed
 - Maintain mappings from physical to virtual memory
 - Decide when to remove a process from memory

- A big chunk of the OS kernel deals with I/O
 - Hundreds of thousands of lines in NT
- The OS provides a standard interface between programs & devices
 - File system (disk), sockets (network), frame buffer (video)
- Device drivers are the routines that interact with specific device types
 - Encapsulates device-specific knowledge
 - e.g., how to initialize a device, request I/O, handle errors
 - Examples: SCSI device drivers, Ethernet card drivers, video card drivers, sound card drivers, ...

Secondary storage

- Secondary storage (disk, tape) is persistent memory
 - Often magnetic media, survives power failures (hopefully)
- Code interacting with disks are at a very low level in the OS
 - Used by many components (file system, VM, ...)
 - Handle scheduling of disk operations, head movement, error handling, and often management of space on disks
- Usually independent of file system
 - Although there may be cooperation
 - FS knowledge of device details can help with performance
 - e.g., place related files close together on disk

File systems

- Storage devices are hard to work with
 - File system offers a convenient abstraction
 - Defines logical abstractions/objects like files & directories
 - As well as operations on these objects
- A file is the basic unit of long-term storage
- A directory is just a special kind of file
 - ... containing names of other files & metadata
- Interface
 - File/directory creation/deletion, manipulation, copy, lock
- Other higher level services: accounting & quotas, backup, indexing or search, versioning

Protection

- Protection is a general mechanism used throughout the OS
 - All resources must be protected
 - memory
 - processes
 - files
 - devices
 - ...
- Mechanisms help detect and contain errors, and preventing malicious destruction

And now a short break ...

Supported Features



OS design & implementation

- A challenge ...
 - Conflicting user and system goals
 - User – Convenient, easy to learn, reliable, safe, fast
 - System – Easy to design, implement, and maintain, flexible, reliable, error-free and efficient
 - Dealing with concurrency – users and devices
 - Some hostile users, others who want to collaborate
 - Long expected lives (Unix is ~40y, Windows ~ 30y) & no clear ideas on future needs
 - Portability and support for 1,000s of devices
 - Backward compatibility

OS design & implementation


- A software engineering principle – separate policy & mechanism
 - Policy: What will be done?
 - Mechanism: How to do it?
 - Why do you care? Max flexibility, easier to change
- Implementation on high-level language
 - Early on – assembly (e.g. MS-DOS – 8088), later Algol (MCP), PL/1 (MULTICS), C (Unix, ...)
 - Advantages – faster to write, more compact, easier to maintain & debug, easier to port
 - Cost – Size, speed?, but who cares?!

OS structure

- OS made of number of components
 - Process, memory management, ...
 - and system programs
 - e.g., bootstrap code, the init program, ...
- Major design issue
 - How do we organize all this?
 - What are the modules, and where do they exist?
 - How do they interact?
http://www.makelinux.net/kernel_map/

Monolithic design

- Major advantage
 - Cost of module interactions is low (procedure call)
- Disadvantages
 - Hard to understand
 - Hard to modify & maintain
 - Unreliable (no isolation between system modules)
- *Alternative?*
 - How to organize the OS in order to simplify design, implementation and maintenance?



Layering – Traditional approach

- Implement OS as a set of layers
 - Each layer presents an enhanced ‘virtual mach’ to the layer above
 - Each can be tested and verified independently

Layer	Function
5	The operator
4	User programs
3	I/O management
2	Operator-process communication
1	Memory and drum management
0	Processor allocation and multiprogramming

Dijkstra's THE system

Problems with layering


- Imposes hierarchical structure
 - but real systems have complex interactions
 - Strict layering isn't flexible enough
- Poor performance
 - Each layer crossing has an associated overhead
- Disjunction between model and reality
 - Systems modelled as layers, but not built that way

HAL – Hardware Abstraction Layer

- An example of layering in modern OSs
- Goal – to hide differences in hardware from most of the OS kernel
 - On a PC, you can consider it as the driver of the motherboard
 - BSD, Mac OS X, Windows NT, Linux, NetBSD all use a HAL either explicitly identified or not

Virtual machines

- Initial release of OS/360 were strictly batch but users wanted timesharing
 - IBM CP/CMS, later renamed VM/370 ('79)
- Timesharing systems provides
 - Multiprogramming & Extended (virtual) machine
- Essence of VM/370 – separate the two
 - Heart of the system (VMM) does multiprogramming & provides multiple exact copies of bare HW to next layer up
 - Each VM can run any OS




Virtual machines

- A resurgence in mid-90s, started with Rosenblum's work on Disco and VMWare
 - Nowadays ... Java VM, Xen, VritulaBox, Virtual Iron, VMLite, Simics, Parallels, Palacios, QEMU, ...
- What for?
 - Server consolidation – from different services in different lightly used machine (administration cost)
 - Different applications for other OS in your desktop
 - Testing and debugging

Microkernels

- Popular in the late 80's, early 90's
 - Recent resurgence
- Goal: Minimal kernel, the rest in user-level
- What for?
 - Better reliability (isolation between components)
 - Ease of extension and customization
 - Poor performance (user/kernel boundary crossings)
- First microkernel Hydra (CMU, 1970)
 - ... Mach (CMU), Chorus (UNIX-like), OS X (Apple), in some ways NT (Microsoft), L4 (Karlsruhe), MINIX 3, ...

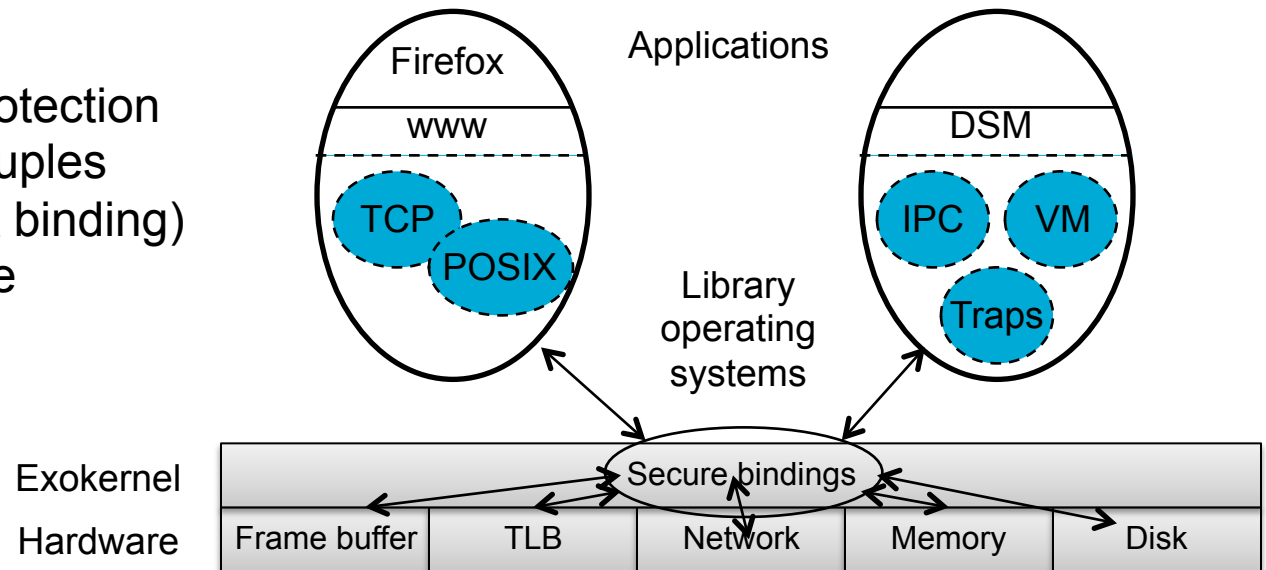
Microkernel



Exokernels

- OS, typically securely multiplexes & *abstract* physical resources
- But no OS abstractions fits all!
- Exokernel
 - A minimal OS securely multiplexes resources
 - Library OSes implement higher-level abstractions

Secure binding – a protection mechanism that decouples authorization (done at binding) from use of a resource



Summary & preview

- Today
 - The mess under the carpet
 - Basic concepts in OS
 - OS design has been an evolutionary process
 - Structuring OS - a few alternatives, not a clear winner
- Next ...
 - Process – the central concept in OS
 - Process model and implementation
 - What it is, what it does and how it does it