

Scheduling I

To do ...

- ❑ Introduction to scheduling
- ❑ Classical algorithms
- ❑ Next Time: Advanced topics on scheduling

Scheduling out there

- You are the manager of a supermarket (ok, things don't always turn out the way we plan them!)
- It's a busy time at 5-6PM and you have one register working; how do you manage the queue to reduce waiting time?
 - You have a handful of customers waiting, each with about equally filled carts
 - A guy, apparently planning to go into hiding, is now in front of the queue and a bunch of people with 2-3 items wait behind
 - An 8-month expectant mother has joined the back of the queue
 - ...

Scheduling

- Problem
 - Several ready processes and much fewer CPUs
- A choice has to be made
 - By the *scheduler*, using a *scheduling algorithm*
- The decision, *scheduling*, is policy
- *Context switching* is a mechanism

Scheduling through time

- Early batch systems – Just run the next job in the tape
- Early timesharing systems – Scarce CPU time, scheduling is critical
- PCs – Commonly one active process, scheduling is easy; with fast and per-user CPU, scheduling is not critical
- Networked workstations, servers, data centers – All back again, multiple ready processes and expensive context switching, scheduling is critical
- Mobile devices, battery life overhead so can the scheduler help
- ...

Environments and goals

- Different scheduling algorithms, with different goals, for different application areas
 - Batch, interactive, real-time
- Batch systems
 - CPU utilization – keep CPU busy (*anything wrong?*)
 - Throughput – max. jobs per hour
 - Turnaround time – min. time bet/ submission & termination
- Interactive systems
 - Response time – handle requests quickly (start responding)
 - Proportionality – meet users' expectations
- ...



Environments and goals

...

- Real-time system
 - Meeting deadlines – avoid losing data
 - Predictability – avoid quality degradation in multimedia
- Average, maximum, minimum or *variance*?
 - Throughput, turnaround time, response time ...
- Goals for all/most systems
 - Fairness – comparable processes getting comparable service
 - Policy enforcement – seeing that stated policy is carried out
 - Balance – keeping all parts of the system busy (mix pool of processes)


Process behavior

- Task – a request to scheduled (multiple tasks/process)
- Workload – set of tasks to run, input to sched algorithm
- Bursts of CPU usage alternate with periods of I/O wait
 - Key to scheduling – CPU-bound & I/O bound process
 - As CPU gets faster – more I/O bound processes



When to schedule?

- When to make scheduling decisions?
 1. At process creation
 2. When a process exits
 3. When a process blocks on I/O, a semaphore, etc
 4. When an I/O interrupts occurs
 5. A fix periods of time – Need a HW clock interrupting




When to schedule?

- A fixed periods of times ... preemptive or not
 - No-preemptive
 - Once a process gets the CPU, it doesn't release it until the process terminates or switches to waiting
 - It may take hours to finish, that's OK
 - Preemptive
 - Let a process run for a maximum fixed time, if running still the OS can preempt the CPU
 - This requires having a clock interrupt at the end of the interval

Comparing policies

- We'll see some *example policies* – in practice, any real system uses some hybrid approach
- For comparison – a metric: turnaround time (performance)

$$T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$$




Process	Burst Time
P1	20
P2	204
P3	8

$$P_1 \text{ turnaround} = 30 - 0 = 30u$$

Comparing policies

- Other metrics/goals (sometimes in conflict)
 - Maximize *CPU utilization*
 - Maximize *throughput* (requests completed / sec)
 - Minimize *average response time* (avg. time from submission of request to first response)
 - Minimize *energy* (joules per instruction) subject to some constraint (e.g., frames/sec)

And now a short break ...




Quino

First-Come First-Served scheduling

- First-Come First-Served (FCFS or FIFO)
 - Simplest, easy to implement, non-preemptive

Process	Burst Time
P1	10
P2	10
P3	10



Avg turnaround time:
 $(10 + 20 + 30)/3 = 20$

Different burst times

Process	Burst Time
P1	24
P2	3
P3	3



Avg turnaround time:
 $(30 + 3 + 6)/3 = 13$

FCFS Issues

Remember the guy going into hiding?

Process	Burst Time
P1	24
P2	3
P3	3



Turnaround time:
 $(24 + 27 + 30)/3 = 27$

- The convoy effect
 - 1 CPU-bound process (burst of 1 sec.)
 - Many I/O-bound ones (needing to read 1000 records)
 - Each I/O-bound process reads one block per sec!
- Potentially bad average response time
- May lead to poor utilization of resources
 - Poor overlap of CPU and I/O

Shortest Job First (SJF)

- Taken from Operation Research



Turnaround time:
 $(30 + 3 + 6)/3 = 13$

Process	Burst Time
P1	24
P2	3
P3	3

- Provably optimal wrt turnaround time

First job finishes at time a; second job at time a + b; ...

Mean turnaround time
 $(4a + 3b + 2c + d)/4$




Biggest
contributor

Job #	Finish time
1	a
2	b
3	c
4	d

Shortest Job First

- Another example



Turnaround time: $(16 + 5 + 1 + 9)/4 = 7.75$

Process	Burst Time
P1	7
P2	4
P3	1
P4	4

- What if they don't all arrive at the same time?*




Turnaround time: $(7 + 11 + 6 + 13)/4 = 9.25$

**Note P2 run at 12 but arrived at 1, so it only waited 11; similar with P3 and P4.*

Process	Arrival	Burst Time
P1	0.0	7
P2	1.0	4
P3	2.0	1
P4	3.0	4

Shortest Remaining Time First

- A preemptive variation



Process	Arrival	Burst Time
P1	0.0	7
P2	1.0	4
P3	2.0	1
P4	3.0	4

Turnaround time: $(16 + 5 + 1 + 7)/4 = 7.25$

- *Great, but how do you know the burst time?*


Determining length of next CPU burst

- Can only *estimate* length
- Typically done using length of previous CPU bursts and exponential averaging

- t_n = actual length of n^{th} CPU burst
- τ_{n+1} = predicted value for the next CPU burst
- $\alpha, 0 \leq \alpha \leq 1$
- Define :

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

Weight of history
 ↓
 Most recent information Past history



$$\tau_1 = 0.5 t_0 + (1 - 0.5) \tau_0 = 8$$

time	0	1	2	3	4	5	6	7	8	9	10
CPU burst (t_i)	6	4	6	4	13	13	13	...			
"guess" (τ_i)	10	8	6	6	5	9	11	12	...		

An arrow points from the equation $\tau_1 = 0.5 t_0 + (1 - 0.5) \tau_0 = 8$ to the value 8 in the "guess" row of the table, which is circled. Another circle highlights the values 6 and 4 in the CPU burst row at time 1.

Priority scheduling

- SJF, a special case of priority-based sched
 - Priority = reverse of predicted next CPU burst
- Pick process with highest priority (lowest #)

Process	Burst time	Priority
P1	10	3
P2	1	1
P3	2	4
P4	1	5
P5	5	2




$$\text{Turnaround time} = (16 + 1 + 18 + 19 + 6)/5 = 12$$

Priority scheduling issues

- How do you assign priorities?
- Possible starvation
 - With an endless supply of high priority jobs, low priority processes may never execute
 - *What other recently discussed algorithm has the same problem?*
 - Solutions
 - Increases priority with age, i.e. accumulated waiting
 - Lower priority as a function of acc'd processing time
 - Assigned maximum quantum

Round-robin scheduling

- SJF is not bad if you know burst times or can estimate it fairly well – the case in many early batch systems
 - At least when measuring turnaround time
- Time-sharing machines changed it all
 - Users want interactivity
 - Turnaround time is not a good metric for this
 - *Response time?* Time to first run minus time of arrival




Turnaround time: $(16 + 5 + 1 + 9)/4 = 7.75$

Response time: $(9 + 1 + 0 + 5)/4 = 3.75$

Process	Burst Time
P1	7
P2	4
P3	1
P4	4

Round-robin scheduling

- Simple, fair, easy to implement, & widely-used
- Each process gets a fix *quantum* or *time slice*
- If quantum expires, preempt CPU
- With n processes & quantum q , each gets $1/n$ of CPU time, no-one waits more than $(n-1) q$ to run first (i.e., response time)




Response time: $(0 + 2 + 4 + 5)/4 = 2.75$

Process	Burst Time
P1	7
P2	4
P3	1
P4	4

Turnaround time: $(16 + 11 + 5 + 13)/4 = 11.25$

Quantum & Turnaround time

- Length of quantum
 - Too short – low CPU efficiency (*why?*)
 - Too long – low response time (*really long, what do you get?*)
 - Commonly ~ 50-100 msec.



Process	Time
P ₁	6
P ₂	3
P ₃	1
P ₄	7

Next time

- How do you support responsive, flexible scheduling? Priority? How are priorities set?
- How do you optimize turnaround time while minimizing response time?
 - Shortest Job First reduces turnaround time but hurts response time
 - Round Robin reduces response time but hurts average waiting time
 - ...?