

Virtual Memory Design and Implementation

To do ...

- ❑ Page replacement algorithms
- ❑ Design and implementation issues
- ❑ Next: Last on virtualization – VMMs

Loading pages

- When should the OS load pages?
 - On demand or ahead of need
- Demand paging
 - Only load when reference
 - Just the code/data needed has to be loaded
 - But needs change over time...
- Anticipatory paging (prefetching)
 - If you know/can guess the pages a process will need
- Few systems try to anticipate future needs
 - The OS is not very good at prediction

And if the page is not there

- If a process references a virtual address in an evicted or never loaded page ...
 - When page was evicted, OS set PTE as invalid and noted disk location of page
 - In a data structure ~page table but holding disk addresses
 - When process tries to access page, page fault
 - OS runs the page fault handler
 - Handler uses the “~PT” data structure to find page on disk
 - ... reads page in, updates PTE to point to it, set it to valid
 - OS restarts the faulting process
 - ... there are a million and one details ...

Page replacement algorithms

- Need room for new page? Page replacement
- What do you do with the victim page?
 - If modified, save it, otherwise just write over it
 - Better not to choose an often used page
- How can any of this work?!?!
 - Locality!

Locality!

- Temporal and spatial locality
 - Temporal – Recently referenced locations tend to be referenced again soon
 - Spatial – Referenced locations tend to be clustered
- Locality means paging could be infrequent
 - A page brought in, gets to be used many times
 - Some issues that may play against this
 - Degree of locality of application
 - Page replacement policy and application reference pattern
 - Amount of physical memory and application footprint

The best page to evict

- Goal of the page replacement algorithm
 - Reduce fault rate by selecting best victim page
 - What's your best candidate for removal?
 - The one you will never touch again – duh!
 - “Never” is a long time
 - For Belady's algorithm – let's say for the longest period
- Let's look at some algorithms
 - for now, assume that a process pages against itself, using a fixed number of page frames

Optimal algorithm (Belady's algorithm)

- Provably optimal
 - Replace page needed at the farthest point in future
 - Optimal but unrealizable
- Estimate by ...
 - Logging page use on previous runs of process
 - Although impractical, useful for simulation

Need room
for this one!

You ideal
victim!

Four page
frames

Reference	A	B	A	C	B	D	A	D	E	D	A	E	B	A	C
1	A		+				+								
2		B			+										
3				C											
4						D		+							

Reference
stream

Compulsory misses
(others: capacity misses)

FIFO algorithm

- Keep a linked list of pages – in order of arrival
- Victim is first page of list
 - Maybe the oldest page will not be used again ...
- Disadvantage
 - But maybe it will – the fact is, you have no idea!
 - Increasing physical memory might increase page faults (Belady's anomaly)

Reference	A	B	A	C	B	D	A	D	E	D	A	E	B	A	C
1	A		+				+		E			+			
2					+						A			+	
3				C									B		
4						D		+		+					C

Need room for this one!


First in

Least recently used (LRU) algorithm

- Pages used recently will be used again soon
 - Throw out page unused for longest time
 - Idea: past experience as a predictor of future needs
 - LRU looks at the past, Belady's wants to look at the future
 - How is LRU different from FIFO?
- Must keep a linked list of pages
 - Most recently used at front, least recently used last
 - Update list with every memory reference!!
 - \$\$\$ in mem. bandwidth, algorithm execution time, etc

Second chance algorithm

- Simple modification of FIFO
 - Avoid removing heavily used pages – check the R bit
- Second chance
 - Pages sorted in FIFO order
 - If page has been used, gets another chance – move it to the end of the list of pages, clear R, update timestamp
 - Page list if fault occurs at time 20, A has R bit set (time is loading time)




Most recently loaded

Oldest page

Clock algorithm

- Second chance is reasonable but inefficient
 - Quit moving pages around – move a pointer?
- Clock – ~Second chance but for implementation
 - Keep all pages in a circular list, as a clock, with the hand pointing to the oldest page
 - When page fault
 - Look at page pointed at by hand
 - If $R = 0$, evict page
 - If $R = 1$, clear R & move hand



Not recently used (NRU) algorithm

- Each page has Reference and Modified bits
 - Set when page is referenced, modified
 - R bit set means recently referenced, so clear it regularly
- Pages are classified

How can this occur?

R	M	Class
0	0	Not referenced, not modified (0,0 → 0)
0	1	Not referenced, modified (0,1 → 1)
1	0	Referenced, but not modified (1,0 → 2)
1	1	Referenced and modified (1,1 → 3)

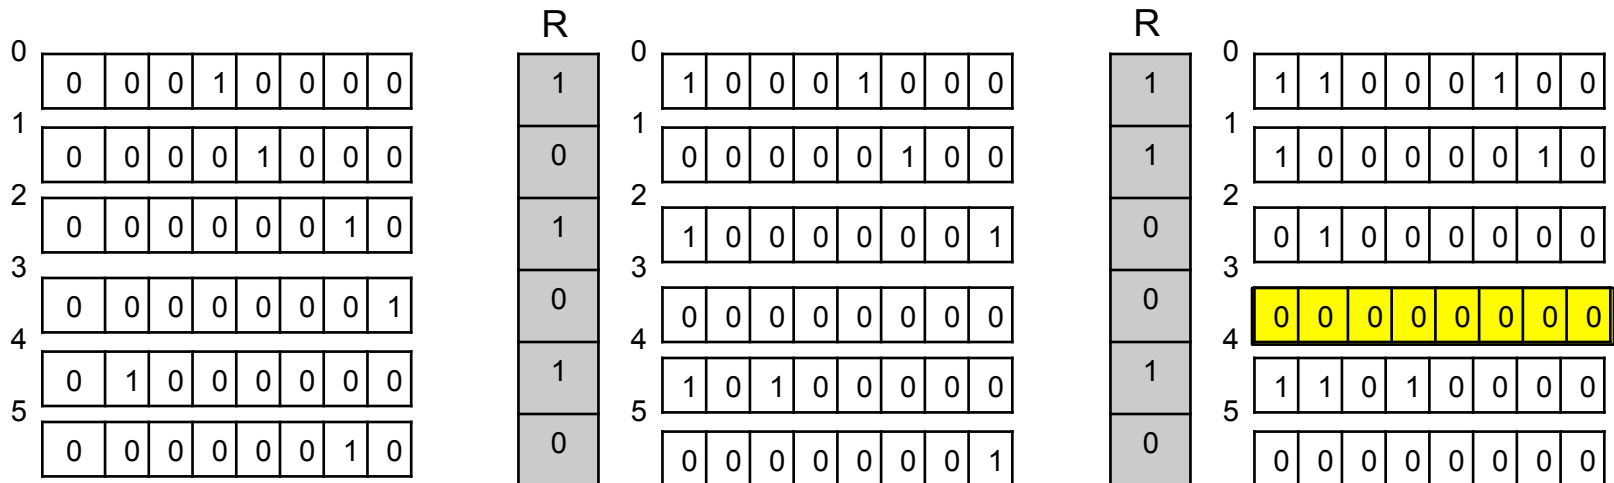
- NRU removes page at random
 - from lowest numbered, non-empty class
- Easy to understand, relatively efficient to implement and sort-of OK performance

Approximating LRU

- With some extra help from hardware
 - Keep a counter in PTE
 - Equipped hardware with a counter, ++ after each instruction
 - After each reference, update PTE counter for the referenced with hardware counter
 - Choose page with lowest value counter
- In software, Not Frequently Used
 - Software counter associated with each page
 - At clock interrupt – add R to counter for each page
 - Problem - it never forgets!
 - A page heavily used early on and never touch again, will keep a high count for a long time and not be evicted

Approximating LRU

- Better, with a small modification – Aging
 - Push R from the left, drop bit on the right
 - How is this not LRU? One bit per tick & a finite number of bits per counter
 - Precision – all pages referenced within the interval are the same
 - Limited past horizon – two pages with all 8 bits unset are the same, even if one dropped a set 9th bit just before




Working set

- Most programs show locality of reference
 - Over a short time, just a few common pages
- Working set
 - Models the dynamic locality of a process' memory usage
 - i.e. the set of pages currently needed by a process
- Intuitively, working set must be in memory, otherwise you'll experience heavy faulting (thrashing)
 - What does it mean 'how much memory does program x need?'
 - What is program x average/worst-case working set size?

Working set

- Demand paging
 - Simplest strategy, load page when needed
 - Can you do better knowing a process WS?
 - Can you use this to reduce turnaround time? Pre-paging
- Working set definition
 - $ws(k, t) = \{p \text{ such that } p \text{ was referenced in the } k \text{ most recent memory references at time } t\}$ (k is WS window size)

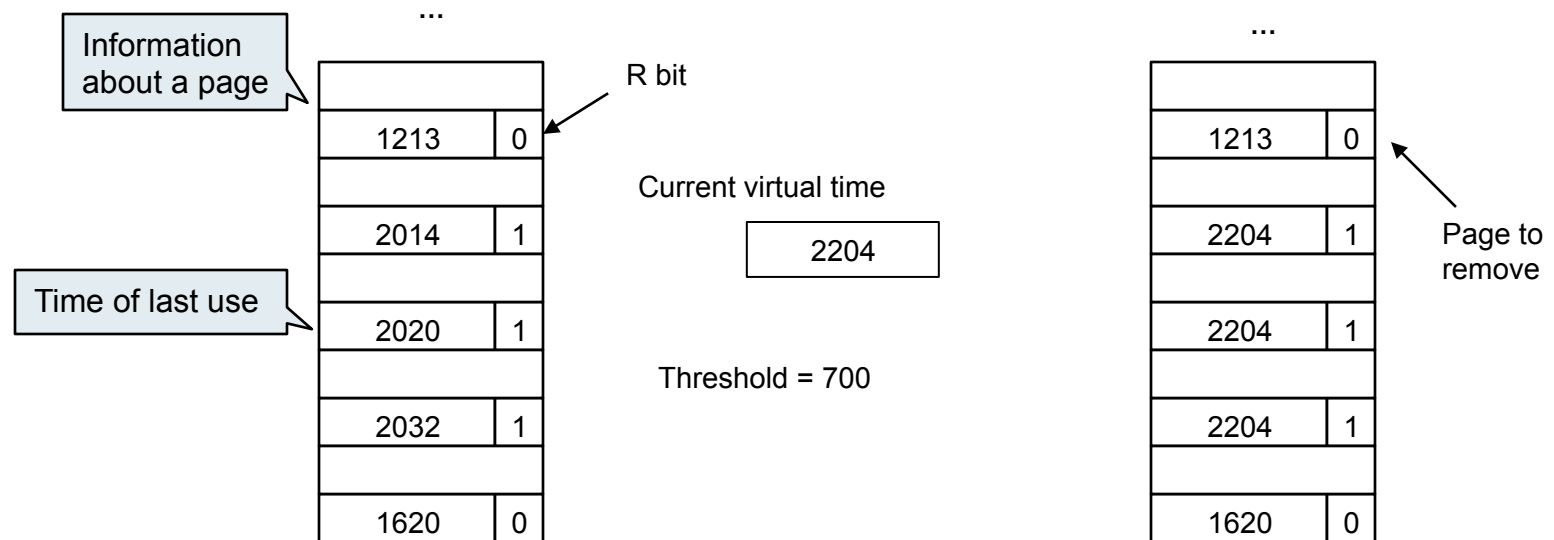
What bounds $ws(k, t)$ as you increase k ?



- A more practical definition – instead of k reference pages, τ msec of execution time (virtual time)


Working set algorithm

- Working set and page replacement
 - Victim – a page not in the working set
- At each clock interrupt – scan the page table
 - $R = 1$? Write Current Virtual Time (CVT) into Time of Last Use
 - $R = 0$? $CVT - \text{Time of Last Use} > \text{Threshold}$? out!
 - Else see if there's some other and evict oldest (w/ $R=0$)
 - If all are in the WS (all $R = 1$), random, preferably clean



WSClock algorithm


- Problem with WS algorithm – Scans the whole table
- Instead, scan only what you need to find a victim
- Combine clock & working set
 - If $R = 1$, unset it
 - If $R = 0$, if $\text{age} > T$ and page clean, out
 - If dirty, schedule write and check next one
 - If loop around,
 - There's 1+ write scheduled
 - you'll have a clean page soon
 - There's none, pick any one



$$R = 0 \text{ \& } 2204 - 1213 > T$$

Cleaning policy

- To avoid having to write pages out when needed – paging daemon
 - Periodically inspects state of memory
 - Keep enough pages free
 - If we need the page before it's overwritten – reclaim it!
- Two hands for better performance (BSD)
 - First one clears R, second checks it
 - If hands are close, only heavily used pages have a chance
 - If back is just ahead of front hand (359°), original clock
 - Two key parameters, adjusted at runtime
 - Scanrate – rate at which hands move through the list
 - Handsread – gap between them



Design issues – global vs. local policy

- When you need a page frame, pick a victim from
 - Among your own resident pages – Local
 - Among all pages – Global
- Local algorithms
 - Basically every process gets a fixed % of memory
- Global algorithms
 - Dynamically allocate frames among processes
 - Better, especially if working set size changes at runtime
 - How many page frames per process?
 - Start with basic set & react to Page Fault Frequency (PFF)
- Most replacement algorithms can work both ways except those based on working set
 - Why not working set based algorithms?

Load control

- Despite good designs, system may still thrash
 - Sum of working sets $>$ physical memory
- How do you know? Page Fault Frequency (PFF)
 - Some processes need more memory
 - But no process needs less ...
- Way out: Swapping
 - So yes, even with paging you still need swapping
 - Reduce number of processes competing for memory
 - ~ two-level scheduling – careful with which process to swap out (there's more than just paging to worry about!)
 - What would you like of the remaining processes?

Backing store

- How do we manage swap area?
 - Allocate space to process when started
 - Keep offset to process swap area in PCB
 - Process can be brought entirely when started or as needed
- Some problems
 - Size – process can grow ... split text/data/stack segments in swap area
 - Do not allocate anything ... you may need extra memory to keep track of pages in swap!

Page fault handling

- Hardware traps to kernel
- General registers saved by assembler routine, OS called
- OS find which virtual page cause the fault
- OS checks address is valid, seeks page frame
- If selected frame is dirty, write it to disk (CS)
- Get new page (CS), update page table
- Back up instruction where interrupted
- Schedule faulting process
- Routine load registers & other state and return to user space

Instruction backup

- With a page fault, the current instruction is stopped part way through ... harder than you think!
 - Consider instruction: MOV.L #6(A1), 2(A0)

One instruction, three memory references (instruction word itself, two offsets for operands

1000

1002

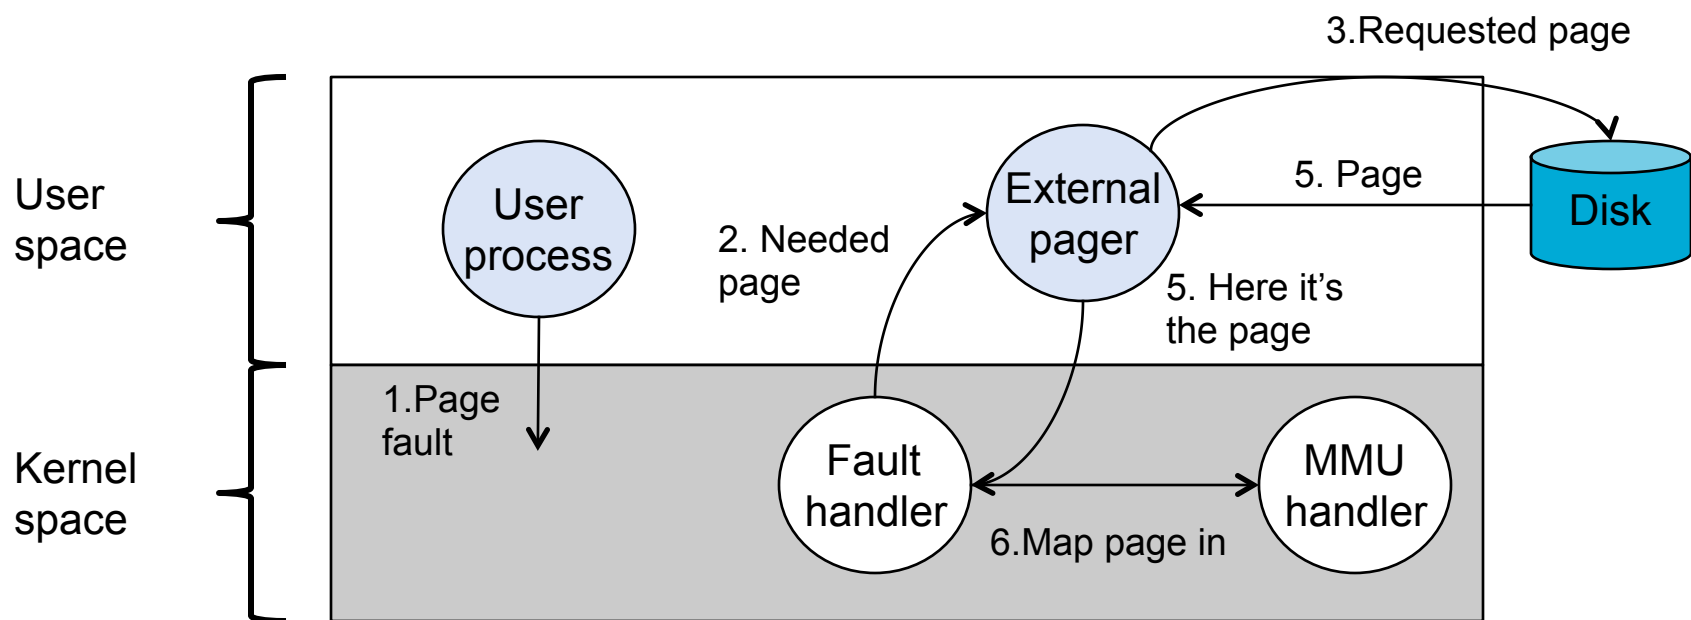
1004

MOVE
6
2

- Which one caused the page fault? What's the PC then?
 - Worse – autodecr/incr as a side-effect of execution?
- Some CPU design include hidden registers to store
 - Beginning of instruction
 - Indicate autodecr./autoincr. and amount

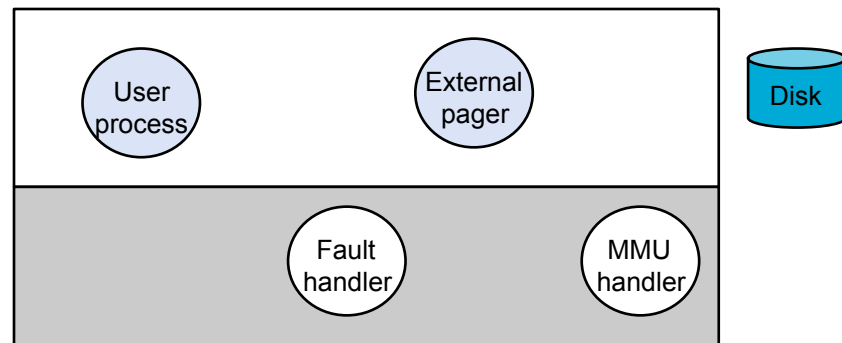
Last: Separating policy & mechanism

- How to structure the memory management system for easy separation? (based on Mach)
 - 1. Low-level MMU handler – machine dependent
 - 2. Page-fault handler in kernel – machine independent, most of paging mechanism
 - 3. External pager running in user space – policy is here



Separation of policy & mechanism

- Where should we put the page replacement algorithm?
 - Cleanest: external pager, but no access to R and M bits
 - Either pass this info to the pager or
 - Fault handler informs external pager which page is victim



- Pros and cons
 - More modular, flexible
 - Overhead of crossing user-kernel + msg exchange
 - As HW get faster and SW more complex ...

Next time

- Virtualize the CPU, virtualize memory, ...
- Let's virtualize the whole machine

And now a short break ...

