

Virtual Machines

To do ...

- ☐ VM over time
- ☐ Implementation methods
- ☐ Hardware features supporting VM
- ☐ Next time: Midterm?

Too many computers!

- An organization can have *too many* machines
 - Why many? Running different services (email, web) to
 - Ensure each has enough resources
 - ... fails independently
 - ... is sandboxed (for security)
 - ... survives if another one gets attack
 - ... different services may require different OSes
 - Why ***too many***? Hard and expensive to maintain!
- Virtualization as an alternative
 - A VMM (or hypervisor) creates the illusion of multiple (virtual) machines

Virtualization applications


- Server consolidation
 - Convert underutilized servers to VMs, saving cost
 - Increasingly used for virtual desktops
- Simplified management
 - Datacenter provisioning and monitoring
 - Dynamic load balancing
- Improved availability
 - Automatic restart, fault tolerance, disaster recovery
- Test and development
- Cloud support
 - Isolation for clients
- *But are you not putting all your eggs in one basket?*

Types of virtualization

- Process virtualization
 - Language-level: Java, .NET, Smalltalk
 - OS-level: processes, Solaris Zones, BSD Jails
 - Cross-ISA emulation: Apple 68K-PPC-x86
- Device virtualization
 - Logical vs. physical: VLAN, VPN, LUN, RAID
- System virtualization
 - Xen, VMware Fusion, KVM, Palacios ...


System virtualization starting point

- Physical hardware
 - Processors, memory, chipset, I/O devices, etc.
 - Resources often grossly underutilized
- Software
 - Tightly coupled to physical hardware
 - Single active OS instance
 - OS controls hardware



Adding a virtualization layer

- Software abstraction
 - Behaves like hardware
 - Encapsulates all OS and application state
- Virtualization layer
 - Extra level of indirection
 - Decouples hardware, OS
 - Enforces isolation
 - Multiplexes physical hardware across VMs



Virtual Machine Monitor

- Classic definition*

... an efficient, isolated duplicate of the real machine. ... the VMM provides an environment ... **essentially identical** with the original machine; second, programs run in this environment show **at worst only minor decreases in speed**; and last, the **VMM is in complete control of system resources**.

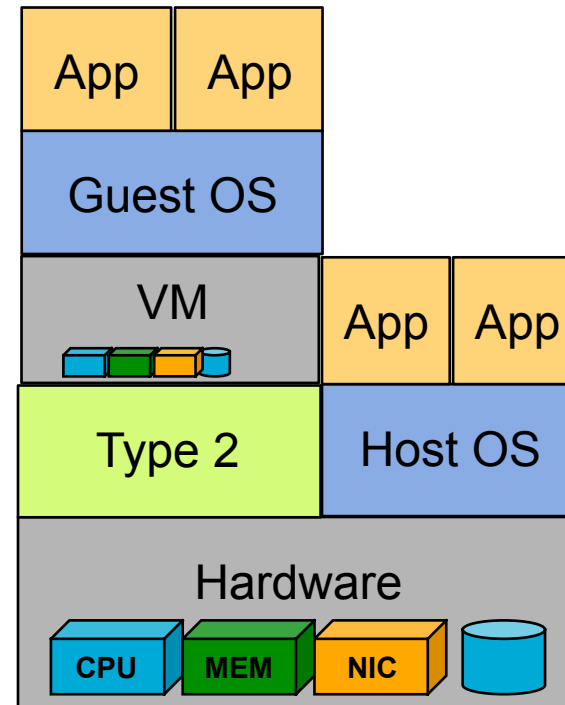
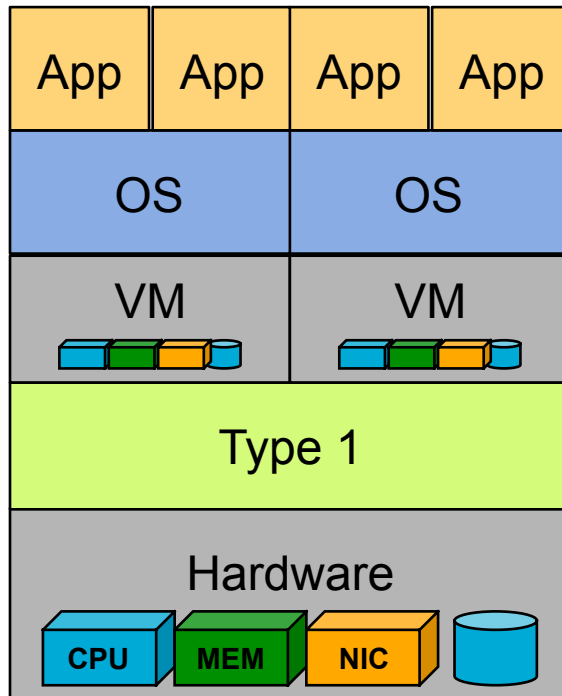
- ➔ VMM properties

- Fidelity – Software on the VMM executes same as on the hardware, other than for timing effects
- Performance – Hardware runs most of the instructions w/o the VMM involvement
- Safety and isolation – VMM manages all hardware resources

*Popek, G. J.; Goldberg, R. P. "Formal requirements for virtualizable third generation architectures". Communications of the ACM, July 1974.


Type 1 and 2 hypervisors

Virtualization method	Type 1	Type 2
Without HW support	ESX Server 1.0	VMware Workstation 1
Paravirtualization	Xen 1.0, Denali	
With HW support	vSphere, Xen, Palacios	VMware Fusion, KVM
Process virtualization		Wine



Classic virtualization

- Classical VMM, back to the 1960s
 - A couple of projects SIMMON and CP-40 for IBM System/360
 - CP-40 → CP-67 → CP/CMS → IBM VM/370
 - IBM mainframes in '72 System/370 (later /390 and z-series)
 - Co-designed proprietary hardware, OS, VMM
 - “Trap and emulate” model
- Applications
 - Timeshare several single-user OS instances on expensive hardware
 - Compatibility





From IBM VM/370 product announcement, *ca.* 1972

Modern virtualization renaissance


- Recent proliferation of VMs
 - Considered exotic mainframe technology in 90s
 - Now pervasive in datacenters and clouds
 - Huge commercial success (partly lead by Rosenblum's VMware, started with DISCO '97)
- Why?
 - Introduction on commodity x86 hardware
 - Ability to “do more with less” saves \$\$\$
 - Innovative new capabilities
 - Extremely versatile technology



And now a short break ...



I'VE GOT A BUNCH OF VIRTUAL WINDOWS MACHINES NETWORKED TOGETHER, HOOKED UP TO AN INCOMING PIPE FROM THE NET. THEY EXECUTE EMAIL ATTACHMENTS, SHARE FILES, AND HAVE NO SECURITY PATCHES.



BETWEEN THEM THEY HAVE PRACTICALLY EVERY VIRUS..

THERE ARE MAILTROJANS, WARHOL WORMS, AND ALL SORTS OF EXOTIC POLYMORPHICS. A MONITORING SYSTEM ADDS AND WIPES MACHINES AT RANDOM. THE DISPLAY SHOWS THE VIRUSES AS THEY MOVE THROUGH THE NETWORK,



GROWING AND STRUGGLING.

YOU KNOW, NORMAL PEOPLE JUST HAVE AQUARIUMS.

GOOD MORNING, BLASTER. ARE YOU AND W32.WELCHIA GETTING ALONG?



WHO'S A GOOD VIRUS? YOU ARE! YES, YOU ARE!


Virtualization's building blocks

- Processor virtualization
 - Trap and Emulate
 - Binary Translation
- Memory virtualization
- I/O virtualization

Virtualizing the CPU

- Running a virtual machine
 - Limited direct execution – remember processes?
 - Want to boot a VM, jump to first address and go ...
- To switch between two VMs, a machine switch
 - Like a processes switch, but using a Virtual CPU
 - VCPU – the state of the CPU as the guest machine believes it to be; a VCPU per guest (~PCB)
 - The VM loaded can be either within the OS itself (i.e., executing a sys call) or within a process running on that OS

- *Easy right?*



Virtualizing the CPU

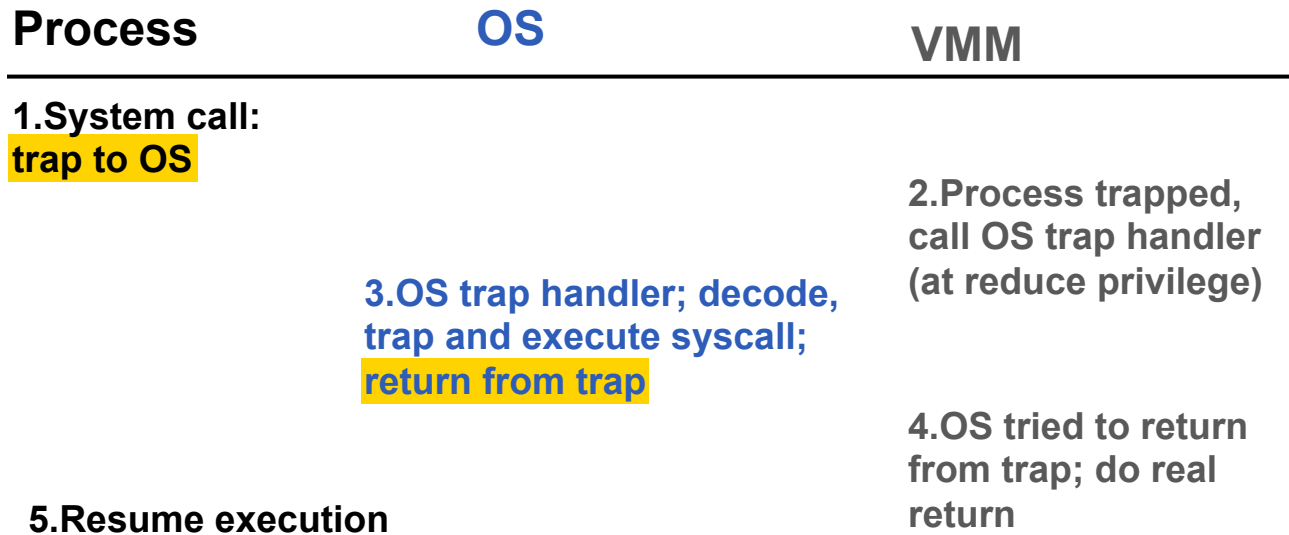
- What if the OS or running app tries to perform a privileged instruction?
 - E.g., update the TLB in a SW managed TLB, or make a syscall
 - The OS cannot be allowed to do it, the VM must intercept it
 - Consider `open(path, flags, mode)` in FreeBSD

```
push dword mode
push dword flags
push dword path
mov eax, 5
push eax
int 80h
```

Process	Hardware	OS
1. Execute instruction (add, load, etc)		
2. System call: trap to OS	3. Switch to kernel mode; jump to trap handler	
		4. In kernel mode; handle syscall; return from trap
	5. Switch to user mode; return to user code	
6. Resume execution		

Trap and emulate

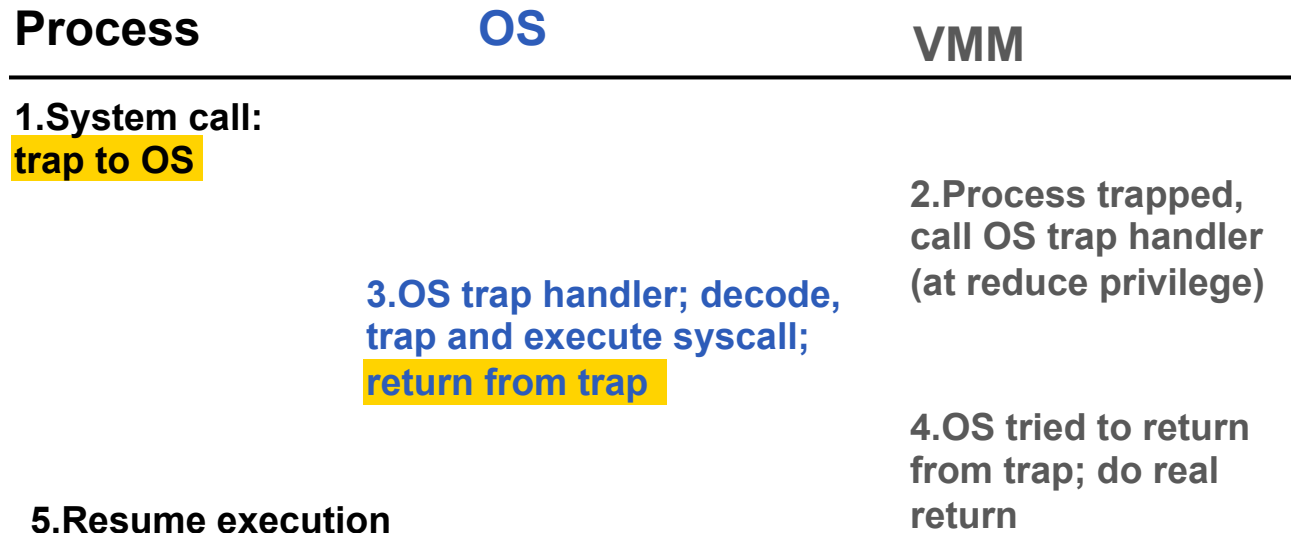
- On a virtualized environment ...trap and emulate
 - Kernel in guest attempts a privileged instruction, traps to VMM
 - VMM emulates the requested action, updates VCPU, and returns control to VM



Why the VMM doesn't handle the system call?

Trap and emulate

- How does the VMM know where the trap handler is?
 - The OS tried to install them at boot time, a privilege instruction, and the VMM took notes



- Trap costs may be high

Trap and emulate

- What mode should the OS run in?
 - Not kernel mode, or it would get unrestricted access to the hardware
- VMM consumes a privilege level
 - In DISCO, use a MIPS supervisor mode
 - Access additional memory, but not privilege instructions
 - Additional memory is enough for the OS to keep its data
 - x86 has four protection mode rings – 3 (least privileged) for user processes, 0 for kernel
 - Without it, run in user mode and virtualize protection levels, use page tables and TLBs to protect OS data structs

Binary translation

- Not all architectures support trap and emulate, i.e. are
- *Strictly virtualizable*
 - A processor is strictly virtualizable if, when executed in a lesser privileged mode:
 - All instructions that access privileged state trap
 - All instructions either trap or execute identically
 - x86 was not strictly virtualizable ☹ until 2005 – example problem instruction – “pop flags” or `popf`
 - In privilege mode, `popf` may change system flags such as IF, which controls interrupt delivery
 - For a de-privileged guest, we need the kernel to trap so that the VMM can emulate the virtual IF
 - *But all user-mode `popf` simply suppresses attempts to modify IF*

Binary translation



- Simple form – fidelity and safety, but bad performance
- Dynamically translate potentially dangerous instructions (non-virtualizable) into safe ones
 - VMM inspects next block of instr (up to a control transfer)
 - Translate each instr and cache translation, jump to start of the translated block and run with VCPU state on HW
- Issues with binary translation
 - Translation cache management
 - PC synchronization on interrupts
 - Self-modifying code
 - Notified on writes to translated guest code
 - Protecting VMM from guest

Virtualization's building blocks

- Processor virtualization
 - Trap and Emulate
 - Binary Translation
- **Memory virtualization**
- I/O virtualization

Virtualizing memory

- Add another level of indirection for memory
- Physical memory is now a virtualization on top of “machine memory”
 - Each OS maps virtual-to-physical addresses via its per-process page table
 - VMM maps physical mappings to underlying machines via its per-OS page tables



Traditional address spaces

- Consider address translation with a software-managed TLB

Process

OS

1. Load from memory;
TLB miss:
trap to OS

2. OS TLB miss handler;
extract VPN from VA;
do page table lookup;
if present and valid,
get PFN, update TLB;
return from trap

3. Resume execution;
instruction is retried,
results in TLB hit

TLB miss flow with virtualization

- With a VMM, upon a miss TLB is not the OS TLB miss handler that runs but the VMM's

Process	OS	VMM
1. Load from memory; TLB miss: trap to OS	3. OS TLB miss handler; extract VPN from VA; do page table lookup; if present and valid, get PFN, update TLB	2. VMM TLB miss handler; call into OS TLB handler (reducing privilege)
	5. Return from trap	4. Trap handler; unprivileged code trying to update TLB; OS is trying to install VPN-PFN mapping; Update TLB instead with VPN-to-MFN (privileged); Jump back to OS (reducing privilege)
6. Resume execution; instruction is retried, results in TLB hit		

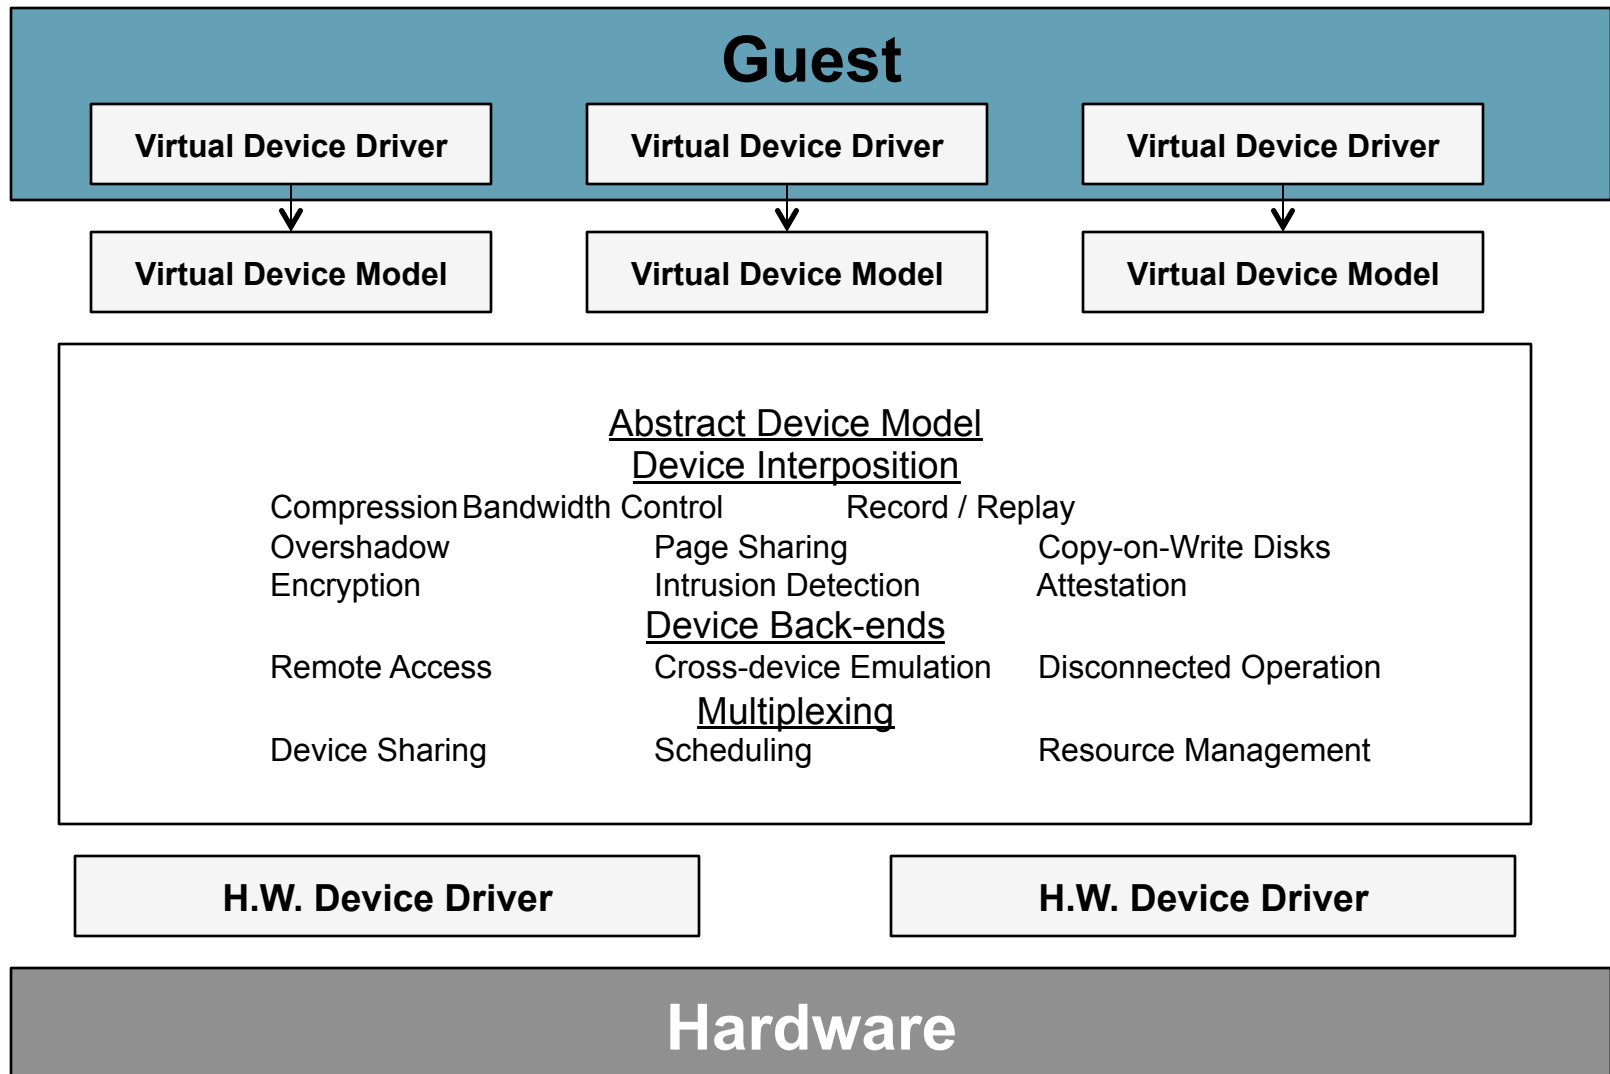
Issues with virtualized memory

- Positives
 - Simplifies monitor design
- TLB misses significantly more expensive
 - DISCO included a software TLB to reduce this cost
- With a hardware-managed TLB
 - HW walks the TLB and updates it as needed
 - VMM must monitor changes the OS makes to each page table and keep a shadow page table mapping the virtual address of each process to the VMM's desired machine pages
 - The VMM installs a process' shadow table when the OS tries to install the process' OS-level page table

Summary

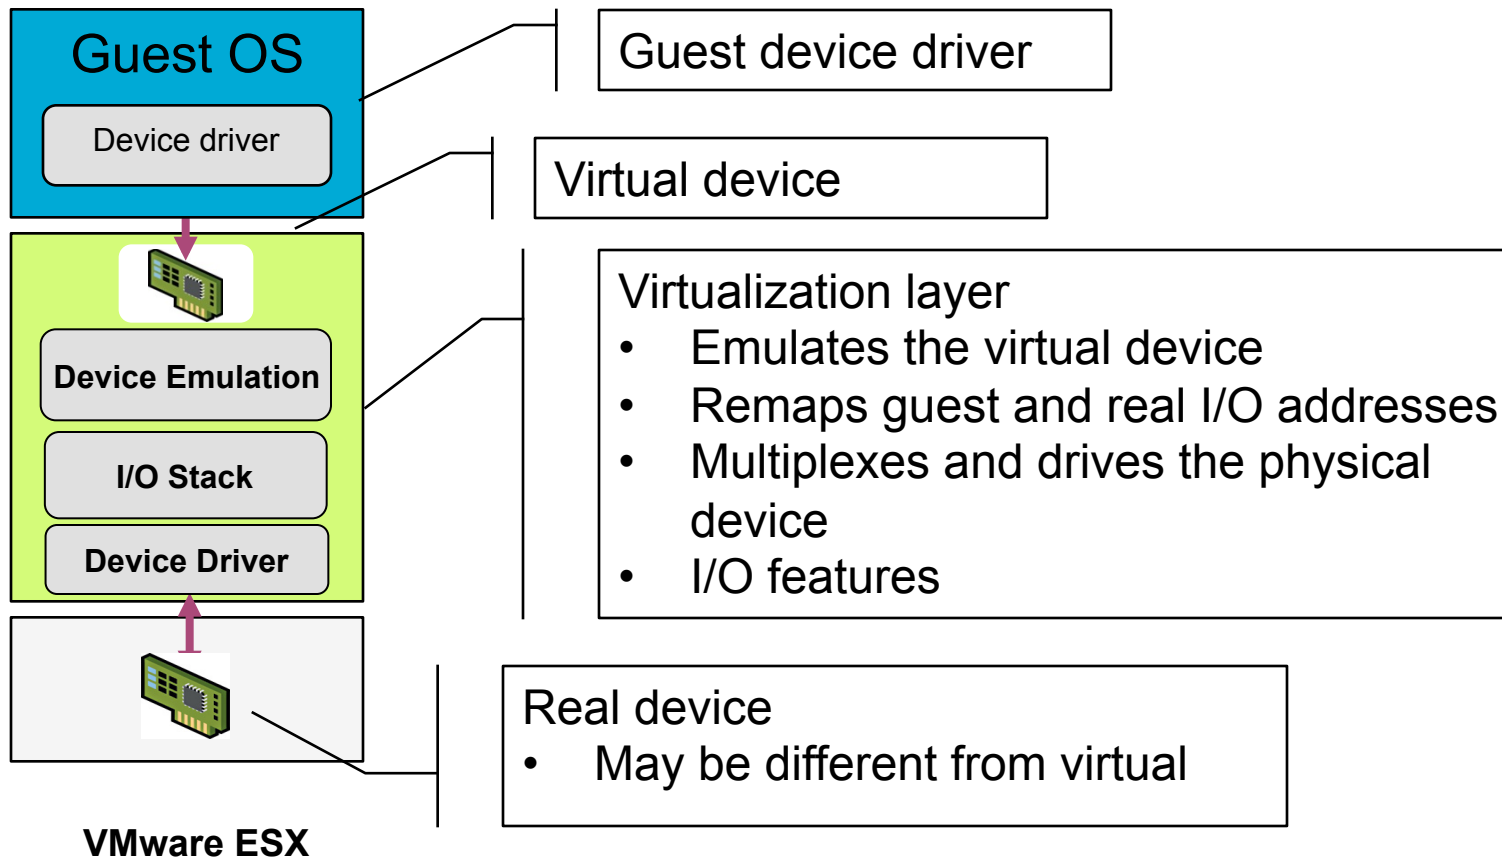
- Virtualization
 - Not a new concept but technological improvements brought a renaissance and many new usages
 - Interested? Take EECS 441 – much of it focused on Linux kernel development
- Second half of the quarter
 - The power and danger of concurrency
 - I/O and persistent storage

I/O Virtualization



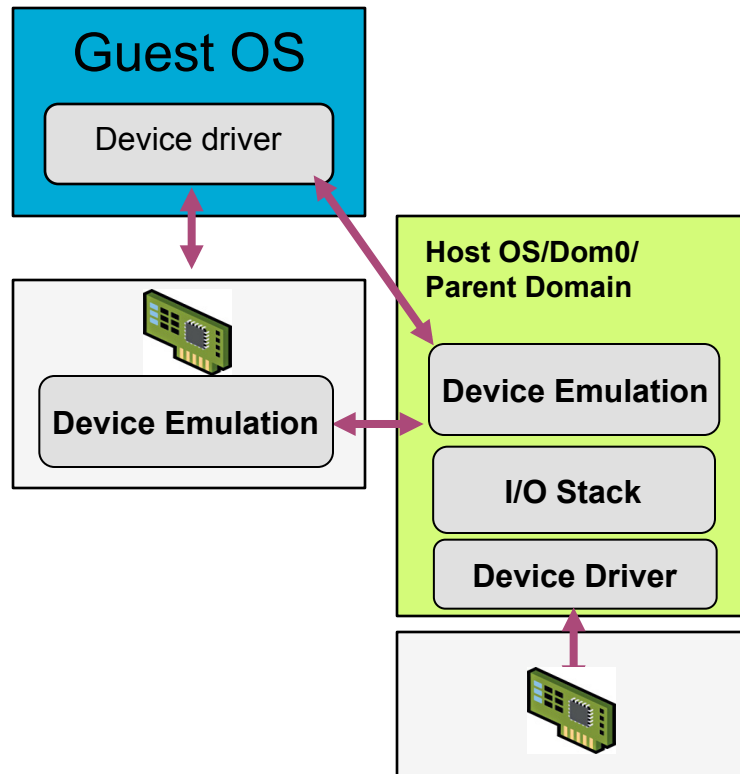
I/O Virtualization implementations

Virtualized I/O



I/O Virtualization implementations

Virtualized I/O: Hosted or Split




VMware Workstation, VMware Server,
Xen, Microsoft Hyper-V, Virtual Server

- Only Domain0, created at boot, has direct access
 - Can also create/terminate domains, control scheduling parameters, physical memory allocation, ..
- Others domains access through virtual device
- Information pass between domains through shared-mem, asynchronous buffer-descriptor rings

I/O Virtualization implementations

Passthrough I/O



VMware ESX (FPT)

- Fast but inflexible
- How to handle
 - Discontinuous physical memory
 - Read-only physical memory (COW)
 - Paged out physical memory
 - VM migration
 - ...